

Exercise 2: Cloud Email Mode with Claude

Experience the difference between local and cloud AI processing

15–20 minutes

Intermediate

CLOUD • Personal Only

Privacy Notice — Read This First

This exercise uses cloud AI models. Your dictation will be sent to external servers for processing. **Do NOT use this mode for PHI, patient data, work emails, or any sensitive information.** This mode is strictly for personal, non-sensitive communications. When in doubt, use your local mode from Exercise 1.

What You'll Build

A cloud-powered email mode using Claude as the AI processor. This exercise isn't just about building another mode—it's about understanding the trade-offs between local and cloud processing so you can make informed decisions about when to use each.

Why Build This If We Already Have a Local Mode?

Because experiencing both approaches teaches you something. Cloud models like Claude are noticeably better at understanding nuance, matching tone, and handling complex instructions. You'll feel the difference when you use it. The point isn't to replace your local mode—it's to understand when cloud processing is worth the privacy trade-off (personal stuff) and when it isn't (work stuff).

Local vs. Cloud: The Trade-Off

Local Mode (Exercise 1)

- ✓ PHI-safe — data never leaves device

Cloud Mode (This Exercise)

- ✗ NOT for sensitive data

- ✓ Works offline
- ✓ Free (no API costs)
- Good formatting
- Basic context understanding

- ✗ Requires internet
- ✗ Requires PRO or API key
- ★ Excellent formatting
- ★ Superior context understanding

The quality difference is real. Claude handles ambiguity better, matches tone more naturally, and produces output that needs less editing. But for government work? Local wins every time because privacy trumps convenience.

Before You Start

- SuperWhisper PRO subscription **OR** your own Anthropic API key
- Exercise 1 completed (you should already understand custom mode basics)
- A personal email account for testing (Gmail, personal Outlook, etc.)

Step-by-Step: Build the Cloud Mode

1

Create a New Custom Mode

Open SuperWhisper Settings → Modes → Add Mode. We're creating a second email mode, so name it clearly.

Name: "Personal Email (Cloud)" or "Smart Email"

The name should remind you this is NOT for work use.

2

Configure Your API Key (If Using Your Own)

If you're using SuperWhisper PRO, cloud models are already available. If you're using your own Anthropic API key, you'll need to add it first.

Go to **Settings** → **AI Models** → **Add Custom Model**

Provider: Anthropic

Model: claude-sonnet-4-20250514 (or latest available)

API Key: [your Anthropic API key]

If you don't have an API key and don't have PRO, you can still read through this exercise to understand the concepts.

3

Select Cloud Models

This time, we're intentionally choosing cloud options for both transcription and AI processing.

Voice Model: Ultra or Nova (CLOUD) — fastest transcription

Language Model: Claude Sonnet (your custom model or PRO built-in)

Yes, this sends audio to servers. That's the point of this exercise—understanding the trade-off.

4

Set Up Auto-Activation (Carefully)

This is where discipline matters. Only activate this mode for personal email—never for work systems.

Click **Add Rule** and add:

For personal Gmail: URL contains "mail.google.com" (but only if this is your personal account)

For personal Outlook: URL contains "outlook.live.com"

DO NOT add outlook.office.com or your work email domain.

Think Before You Add Rules

If your work email and personal email are both in the browser, be very careful with URL patterns. It's better to manually activate this mode than to accidentally send work content to cloud servers. When in doubt, don't add auto-activation at all—just use a keyboard shortcut.

5

Enable All Context Types

Cloud models are better at using context, so let's give Claude everything.

- Application Context** — Knows what app/browser you're in
- Selected Text Context** — Can read the email you're replying to
- Clipboard Context** — Can reference recently copied text

More context = smarter output. Claude handles this well.

6

Enter Advanced AI Instructions

This prompt takes advantage of Claude's superior reasoning. It's more sophisticated than our local mode prompt because Claude can handle the complexity.

```
# Smart Email Mode - Claude-Powered
```

```
<role>
```

```
You are an expert communication assistant specializing in email composition.  
You understand nuance, tone, and social dynamics in professional communication.
```

```
</role>
```

```
<instructions>
```

```
Transform dictated speech into a polished, contextually appropriate email.
```

Before writing, analyze the available context:

- If Selected Text contains an email, this is a REPLY. Match the sender's tone and formality level. Reference specific points from their message.
- If Clipboard contains relevant text, incorporate it naturally.
- Adjust formality based on the detected relationship (close colleague, acquaintance, external contact, formal business).

Structure the email:

```
1. Greeting that matches the relationship and context
2. Clear opening that establishes purpose
3. Body with well-organized points (use bullets if listing multiple items)
4. Specific next steps or call-to-action when appropriate
5. Sign-off that matches the overall tone
</instructions>

<style>
- Match tone to context (casual colleague → casual; formal external → formal)
- Keep sentences concise and scannable
- Use paragraph breaks for readability
- Be specific rather than vague
- Sound like a competent professional, not a robot
</style>

<constraints>
- Never include placeholder text like [Name] or [Date] - infer from context
- If something is ambiguous, make a reasonable assumption and move on
- Output ONLY the email text, ready to paste and send
- No explanations, no meta-commentary, no "Here's your email:"
</constraints>
```

Notice What's Different

This prompt asks Claude to *analyze context* and *make inferences*—things that local models struggle with. We're taking advantage of Claude's strengths: understanding nuance, handling ambiguity, and adapting to social context. You couldn't get this level of sophistication from a local model.

7

Save and Prepare to Test

Save your mode. Before testing, double-check that you're about to use a personal email account—not your work email.

Test Your Cloud Mode

Let's see what Claude can do. Run these tests in your **personal** email only.

Test 1: Fresh Email with Complex Request

1. Open your personal Gmail or email app
2. Start a new email
3. Hold your recording hotkey and dictate something complicated:

"I need to schedule a get-together with the book club people, probably next Saturday or the one after, and I should mention that Sarah suggested we do a potluck this time instead of just ordering pizza again, and also ask if anyone has started the new book yet because I'm only on chapter three and don't want to be behind"

4. Release and watch Claude organize that mess into a coherent email

Test 2: Reply with Tone Matching

1. Find a casual email from a friend or family member
2. Select/highlight the text of their email
3. Hold your recording hotkey and dictate:

"Tell them that sounds great, I'm definitely in, and ask what I should bring"

4. Notice how Claude matches the casual tone of the original email

Test 3: Compare with Local Mode

1. Switch to your local email mode (Exercise 1)
2. Dictate the same Test 1 content
3. Compare the outputs side by side

The local mode output is probably fine—but the Claude output is likely more natural, better organized, and requires less editing. That's the trade-off in action.

□ Push the Limits

Try dictating something deliberately vague: "Reply to that and say I'll think about it but I'm not sure yet." With no context about what "that" refers to, Claude will make reasonable assumptions. Local models often get confused here.

Also try mixing topics in one dictation: scheduling something, asking a question, AND mentioning an unrelated thing. See how Claude organizes multi-topic content versus the local model.

The Workflow

Cloud Mode in Action

Open Personal Gmail



Mode Auto-Activates



Select Email to Reply



Dictate Response



Polished Email Appears

□ THE TWO-MODE STRATEGY

Now you have two email modes: local for work, cloud for personal. This isn't redundant—it's intentional. Use keyboard shortcuts to switch: maybe ⌘+Shift+E for local (work) and ⌘+Shift+P for personal (cloud). Different tools for different jobs.

Completion Checklist

- Created mode named something like "Personal Email (Cloud)"
- Configured Anthropic API key or verified PRO access
- Selected Claude Sonnet as Language Model
- Set auto-activation for personal email ONLY (not work)
- Enabled all three context types
- Pasted the advanced Claude prompt
- Tested with complex fresh email
- Tested reply with tone matching
- Compared output quality with local mode

The Golden Rule

Local mode for work. Cloud mode for personal. Write this on a sticky note if you need to. The quality difference isn't worth a compliance issue. Your local mode is plenty good for professional email—save the cloud power for personal stuff where privacy constraints don't apply.

What's Next?

You now understand the spectrum from local to cloud processing. In **Exercise 3**, we'll explore prompt enrichment—techniques to make your AI instructions more powerful, regardless of which model you're using. These techniques work for both local and cloud modes.